

Papers

The following are the latest papers (sorted by release date) on prompt engineering for large language models (LLMs). We update the list of papers on a daily/weekly basis.

Overviews

- [The Prompt Report: A Systematic Survey of Prompting Techniques](#) (June 2024)
- [Prompt Design and Engineering: Introduction and Advanced Methods](#) (January 2024)
- [A Survey on Hallucination in Large Language Models: Principles, Taxonomy, Challenges, and Open Questions](#) (November 2023)
- [An RL Perspective on RLHF, Prompting, and Beyond](#) (October 2023)
- [Few-shot Fine-tuning vs. In-context Learning: A Fair Comparison and Evaluation](#) (May 2023)
- [Jailbreaking ChatGPT via Prompt Engineering: An Empirical Study](#) (May 2023)
- [Harnessing the Power of LLMs in Practice: A Survey on ChatGPT and Beyond](#) (April 2023)
- [Tool Learning with Foundation Models](#) (April 2023)
- [One Small Step for Generative AI, One Giant Leap for AGI: A Complete Survey on ChatGPT in AIGC Era](#) (April 2023)
- [A Bibliometric Review of Large Language Models Research from 2017 to 2023](#) (April 2023)
- [A Survey of Large Language Models](#) (April 2023)
- [Nature Language Reasoning, A Survey](#) (March 2023)
- [Augmented Language Models: a Survey](#) (February 2023)
- [A Survey for In-context Learning](#) (December 2022)
- [Towards Reasoning in Large Language Models: A Survey](#) (December 2022)
- [Reasoning with Language Model Prompting: A Survey](#) (December 2022)
- [Emergent Abilities of Large Language Models](#) (June 2022)
- [A Taxonomy of Prompt Modifiers for Text-To-Image Generation](#) (April 2022)
- [Pre-train, Prompt, and Predict: A Systematic Survey of Prompting Methods in Natural Language Processing](#) (July 2021)

Approaches

- [Enhancing Zero-Shot Chain-of-Thought Reasoning in Large Language Models through Logic](#) (February 2024)
- [Principled Instructions Are All You Need for Questioning LLaMA-1/2, GPT-3.5/4](#) (December 2023)
- [Walking Down the Memory Maze: Beyond Context Limit through Interactive Reading](#) (October 2023)
- [Large Language Models as Analogical Reasoners](#) (October 2023)
- [LLMLingua: Compressing Prompts for Accelerated Inference of Large Language Models](#) (October 2023)
- [Query-Dependent Prompt Evaluation and Optimization with Offline Inverse RL](#) (September 2023)
- [Chain-of-Verification Reduces Hallucination in Large Language Models](#) (September 2023)
- [Connecting Large Language Models with Evolutionary Algorithms Yields Powerful Prompt Optimizers](#) (September 2023)
- [From Sparse to Dense: GPT-4 Summarization with Chain of Density Prompting](#) (September 2023)
- [Re-Reading Improves Reasoning in Language Models](#) (September 2023)
- [Graph of Thoughts: Solving Elaborate Problems with Large Language Models](#) (August 2023)
- [Skeleton-of-Thought: Large Language Models Can Do Parallel Decoding](#) (July 2023)
- [Focused Prefix Tuning for Controllable Text Generation](#) (June 2023)
- [Exploring Lottery Prompts for Pre-trained Language Models](#) (May 2023)
- [Less Likely Brainstorming: Using Language Models to Generate Alternative Hypotheses](#) (May 2023)
- [Let's Verify Step by Step](#) (May 2023)
- [Universality and Limitations of Prompt Tuning](#) (May 2023)
- [MultiTool-CoT: GPT-3 Can Use Multiple External Tools with Chain of Thought Prompting](#) (May 2023)
- [PEARL: Prompting Large Language Models to Plan and Execute Actions Over Long Documents](#) (May 2023)
- [Reasoning with Language Model is Planning with World Model](#) (May 2023)
- [Self-Critique Prompting with Large Language Models for Inductive Instructions](#) (May 2023)
- [Better Zero-Shot Reasoning with Self-Adaptive Prompting](#) (May 2023)

- [Hierarchical Prompting Assists Large Language Model on Web Navigation](#) (May 2023)
- [Interactive Natural Language Processing](#) (May 2023)
- [Can We Edit Factual Knowledge by In-Context Learning?](#) (May 2023)
- [In-Context Learning of Large Language Models Explained as Kernel Regression](#) (May 2023)
- [Plan-and-Solve Prompting: Improving Zero-Shot Chain-of-Thought Reasoning by Large Language Models](#) (May 2023)
- [Meta-in-context learning in large language models](#) (May 2023)
- [Let's Sample Step by Step: Adaptive-Consistency for Efficient Reasoning with LLMs](#) (May 2023)
- [Post Hoc Explanations of Language Models Can Improve Language Models](#) (May 2023)
- [Compress, Then Prompt: Improving Accuracy-Efficiency Trade-off of LLM Inference with Transferable Prompt](#) (May 2023)
- [TreePrompt: Learning to Compose Tree Prompts for Explainable Visual Grounding](#) (May 2023)
- [TELeR: A General Taxonomy of LLM Prompts for Benchmarking Complex Tasks](#) (May 2023)
- [Efficient Prompting via Dynamic In-Context Learning](#) (May 2023)
- [The Web Can Be Your Oyster for Improving Large Language Models](#) (May 2023)
- [Flatness-Aware Prompt Selection Improves Accuracy and Sample Efficiency](#) (May 2023)
- [Tree of Thoughts: Deliberate Problem Solving with Large Language Models](#) (May 2023)
- [ZeroPrompt: Streaming Acoustic Encoders are Zero-Shot Masked LMs](#) (May 2023)
- [Chain-of-Symbol Prompting Elicits Planning in Large Language Models](#) (May 2023)
- [Cook: Empowering General-Purpose Language Models with Modular and Collaborative Knowledge](#) (May 2023)
- [What In-Context Learning "Learns" In-Context: Disentangling Task Recognition and Task Learning](#) (May 2023)
- [Reprompting: Automated Chain-of-Thought Prompt Inference Through Gibbs Sampling](#) (May 2023)
- [Satisfiability-Aided Language Models Using Declarative Prompting](#) (May 2023)
- [Pre-Training to Learn in Context](#) (May 2023)
- [Boosted Prompt Ensembles for Large Language Models](#) (April 2023)
- [Global Prompt Cell: A Portable Control Module for Effective Prompt](#) (April 2023)

- [Why think step-by-step? Reasoning emerges from the locality of experience](#) (April 2023)
- [Revisiting Automated Prompting: Are We Actually Doing Better?](#) (April 2023)
- [REFINER: Reasoning Feedback on Intermediate Representations](#) (April 2023)
- [Reflexion: an autonomous agent with dynamic memory and self-reflection](#) (March 2023)
- [CAMEL: Communicative Agents for "Mind" Exploration of Large Scale Language Model Society](#) (March 2023)
- [Self-Refine: Iterative Refinement with Self-Feedback](#) (March 2023)
- [kNN Prompting: Beyond-Context Learning with Calibration-Free Nearest Neighbor Inference](#) (March 2023)
- [Visual-Language Prompt Tuning with Knowledge-guided Context Optimization](#) (March 2023)
- [Fairness-guided Few-shot Prompting for Large Language Models](#) (March 2023)
- [Context-faithful Prompting for Large Language Models](#) (March 2023)
- [Is Prompt All You Need? No. A Comprehensive and Broader View of Instruction Learning](#) (March 2023)
- [UPRISE: Universal Prompt Retrieval for Improving Zero-Shot Evaluation](#) (March 2023)
- [Model-tuning Via Prompts Makes NLP Models Adversarially Robust](#) (March 2023)
- [Structure Pretraining and Prompt Tuning for Knowledge Graph Transfer](#) (March 2023)
- [CoTEVer: Chain of Thought Prompting Annotation Toolkit for Explanation Verification](#) (March 2023)
- [Larger language models do in-context learning differently](#) (March 2023)
- [OpenICL: An Open-Source Framework for In-context Learning](#) (March 2023)
- [Dynamic Prompting: A Unified Framework for Prompt Tuning](#) (March 2023)
- [ART: Automatic multi-step reasoning and tool-use for large language models](#) (March 2023)
- [Multitask Prompt Tuning Enables Parameter-Efficient Transfer Learning](#) (March 2023)
- [Effectiveness of Data Augmentation for Prefix Tuning with Limited Data](#) (March 2023)
- [Mixture of Soft Prompts for Controllable Data Generation](#) (March 2023)
- [Prompt, Generate, then Cache: Cascade of Foundation Models makes Strong Few-shot Learners](#) (March 2023)
- [How Robust is GPT-3.5 to Predecessors? A Comprehensive Study on Language Understanding Tasks](#) (March 2023)

- [Can ChatGPT Understand Too? A Comparative Study on ChatGPT and Fine-tuned BERT](#) (February 2023)
- [EvoPrompting: Language Models for Code-Level Neural Architecture Search](#) (February 2023)
- [In-Context Instruction Learning](#) (February 2023)
- [Chain of Hindsight Aligns Language Models with Feedback](#) (February 2023)
- [Language Is Not All You Need: Aligning Perception with Language Models](#) (February 2023)
- [Automatic Prompt Augmentation and Selection with Chain-of-Thought from Labeled Data](#) (February 2023)
- [Active Prompting with Chain-of-Thought for Large Language Models](#) (February 2023)
- [More than you've asked for: A Comprehensive Analysis of Novel Prompt Injection Threats to Application-Integrated Large Language Models](#) (February 2023)
- [A Prompt Pattern Catalog to Enhance Prompt Engineering with ChatGPT](#) (February 2023)
- [Guiding Large Language Models via Directional Stimulus Prompting](#) (February 2023)
- [How Does In-Context Learning Help Prompt Tuning?](#) (February 2023)
- [Scalable Prompt Generation for Semi-supervised Learning with Language Models](#) (February 2023)
- [Bounding the Capabilities of Large Language Models in Open Text Generation with Prompt Constraints](#) (February 2023)
- [À-la-carte Prompt Tuning \(APT\): Combining Distinct Data Via Composable Prompting](#) (February 2023)
- [GraphPrompt: Unifying Pre-Training and Downstream Tasks for Graph Neural Networks](#) (February 2023)
- [The Capacity for Moral Self-Correction in Large Language Models](#) (February 2023)
- [SwitchPrompt: Learning Domain-Specific Gated Soft Prompts for Classification in Low-Resource Domains](#) (February 2023)
- [Evaluating the Robustness of Discrete Prompts](#) (February 2023)
- [Compositional Exemplars for In-context Learning](#) (February 2023)
- [Hard Prompts Made Easy: Gradient-Based Discrete Optimization for Prompt Tuning and Discovery](#) (February 2023)
- [Multimodal Chain-of-Thought Reasoning in Language Models](#) (February 2023)
- [Large Language Models Can Be Easily Distracted by Irrelevant Context](#) (February 2023)
- [Synthetic Prompting: Generating Chain-of-Thought Demonstrations for Large Language Models](#) (February 2023)
- [Progressive Prompts: Continual Learning for Language Models](#) (January 2023)

- [Batch Prompting: Efficient Inference with LLM APIs](#) (January 2023)
- [Demonstrate-Search-Predict: Composing retrieval and language models for knowledge-intensive NLP](#) (December 2022)
- [On Second Thought, Let's Not Think Step by Step! Bias and Toxicity in Zero-Shot Reasoning](#) (December 2022)
- [Constitutional AI: Harmlessness from AI Feedback](#) (December 2022)
- [Successive Prompting for Decomposing Complex Questions](#) (December 2022)
- [Large Language Models are reasoners with Self-Verification](#) (December 2022)
- [Discovering Language Model Behaviors with Model-Written Evaluations](#) (December 2022)
- [Structured Prompting: Scaling In-Context Learning to 1,000 Examples](#) (December 2022)
- [PAL: Program-aided Language Models](#) (November 2022)
- [Large Language Models Are Human-Level Prompt Engineers](#) (November 2022)
- [Ignore Previous Prompt: Attack Techniques For Language Models](#) (November 2022)
- [Machine Generated Text: A Comprehensive Survey of Threat Models and Detection Methods](#) (November 2022)
- [Teaching Algorithmic Reasoning via In-context Learning](#) (November 2022)
- [Enhancing Self-Consistency and Performance of Pre-Trained Language Models through Natural Language Inference](#) (November 2022)
- [Ask Me Anything: A simple strategy for prompting language models](#) (October 2022)
- [Recitation-Augmented Language Models](#) (October 2022)
- [ReAct: Synergizing Reasoning and Acting in Language Models](#) (October 2022)
- [Prompting GPT-3 To Be Reliable](#) (October 2022)
- [Decomposed Prompting: A Modular Approach for Solving Complex Tasks](#) (October 2022)
- [Automatic Chain of Thought Prompting in Large Language Models](#) (October 2022)
- [Language Models Are Greedy Reasoners: A Systematic Formal Analysis of Chain-of-Thought](#) (October 2022)
- [Evaluating the Susceptibility of Pre-Trained Language Models via Handcrafted Adversarial Examples](#) (September 2022)
- [Dynamic Prompt Learning via Policy Gradient for Semi-structured Mathematical Reasoning](#) (September 2022)
- [Promptagator: Few-shot Dense Retrieval From 8 Examples](#) (September 2022)
- [Atlas: Few-shot Learning with Retrieval Augmented Language Models](#) (November 2022)
- [DocPrompting: Generating Code by Retrieving the Docs](#) (July 2022)
- [On the Advance of Making Language Models Better Reasoners](#) (June 2022)

- [Large Language Models are Zero-Shot Reasoners](#) (May 2022)
- [Maieutic Prompting: Logically Consistent Reasoning with Recursive Explanations](#) (May 2022)
- [MRKL Systems: A modular, neuro-symbolic architecture that combines large language models, external knowledge sources and discrete reasoning](#) (May 2022)
- [PPT: Pre-trained Prompt Tuning for Few-shot Learning](#) (May 2022)
- [Toxicity Detection with Generative Prompt-based Inference](#) (May 2022)
- [Learning to Transfer Prompts for Text Generation](#) (May 2022)
- [The Unreliability of Explanations in Few-shot Prompting for Textual Reasoning](#) (May 2022)
- [A Taxonomy of Prompt Modifiers for Text-To-Image Generation](#) (April 2022)
- [PromptChainer: Chaining Large Language Model Prompts through Visual Programming](#) (March 2022)
- [Self-Consistency Improves Chain of Thought Reasoning in Language Models](#) (March 2022)
- [Training language models to follow instructions with human feedback](#)
- [Rethinking the Role of Demonstrations: What Makes In-Context Learning Work?](#) (February 2022)
- [Chain of Thought Prompting Elicits Reasoning in Large Language Models](#) (January 2022)
- [Show Your Work: Scratchpads for Intermediate Computation with Language Models](#) (November 2021)
- [AI Chains: Transparent and Controllable Human-AI Interaction by Chaining Large Language Model Prompts](#) (October 2021)
- [Generated Knowledge Prompting for Commonsense Reasoning](#) (October 2021)
- [Multitask Prompted Training Enables Zero-Shot Task Generalization](#) (October 2021)
- [Reframing Instructional Prompts to GPTk's Language](#) (September 2021)
- [Design Guidelines for Prompt Engineering Text-to-Image Generative Models](#) (September 2021)
- [Making Pre-trained Language Models Better Few-shot Learners](#) (August 2021)
- [Fantastically Ordered Prompts and Where to Find Them: Overcoming Few-Shot Prompt Order Sensitivity](#) (April 2021)
- [BERTese: Learning to Speak to BERT](#) (April 2021)
- [The Power of Scale for Parameter-Efficient Prompt Tuning](#) (April 2021)
- [Prompt Programming for Large Language Models: Beyond the Few-Shot Paradigm](#) (February 2021)
- [Calibrate Before Use: Improving Few-Shot Performance of Language Models](#) (February 2021)

- [Prefix-Tuning: Optimizing Continuous Prompts for Generation](#) (January 2021)
- [Learning to Generate Task-Specific Adapters from Task Description](#) (January 2021)
- [Making Pre-trained Language Models Better Few-shot Learners](#) (December 2020)
- [Learning from Task Descriptions](#) (November 2020)
- [AutoPrompt: Eliciting Knowledge from Language Models with Automatically Generated Prompts](#) (October 2020)
- [Language Models are Few-Shot Learners](#) (May 2020)
- [How Can We Know What Language Models Know?](#) (July 2020)
- [Scaling Laws for Neural Language Models](#) (January 2020)

Applications

- [PromptRE: Weakly-Supervised Document-Level Relation Extraction via Prompting-Based Data Programming](#) (October 2023)
- [Prompting Large Language Models with Chain-of-Thought for Few-Shot Knowledge Base Question Generation](#) (October 2023)
- [Who Wrote it and Why? Prompting Large-Language Models for Authorship Verification](#) (October 2023)
- [Promptor: A Conversational and Autonomous Prompt Generation Agent for Intelligent Text Entry Techniques](#) (October 2023)
- [Thought Propagation: An Analogical Approach to Complex Reasoning with Large Language Models](#) (October 2023)
- [From Sparse to Dense: GPT-4 Summarization with Chain of Density Prompting](#) (September 2023)
- [Self-Taught Optimizer \(STOP\): Recursively Self-Improving Code Generation](#) (October 2023)
- [Think before you speak: Training Language Models With Pause Tokens](#) (October 2023)
- [\(Dynamic\) Prompting might be all you need to repair Compressed LLMs](#) (October 2023)
- [In-Context Learning in Large Language Models: A Neuroscience-inspired Analysis of Representations](#) (September 2023)
- [Understanding In-Context Learning from Repetitions](#) (September 2023)
- [Investigating the Efficacy of Large Language Models in Reflective Assessment Methods through Chain of Thoughts Prompting](#) (September 2023)
- [Automatic Prompt Rewriting for Personalized Text Generation](#) (September 2023)
- [Efficient Streaming Language Models with Attention Sinks](#) (September 2023)
- [The Dawn of LMMs: Preliminary Explorations with GPT-4V\(ision\)](#) (September 2023)

- [Graph Neural Prompting with Large Language Models](#) (September 2023)
- [Large Language Model Alignment: A Survey](#) (September 2023)
- [Enhancing Zero-Shot Chain-of-Thought Reasoning in Large Language Models through Logic](#) (September 2023)
- [A Practical Survey on Zero-shot Prompt Design for In-context Learning](#) (September 2023)
- [EchoPrompt: Instructing the Model to Rephrase Queries for Improved In-context Learning](#) (September 2023)
- [Prompt, Condition, and Generate: Classification of Unsupported Claims with In-Context Learning](#) (September 2023)
- [PolicyGPT: Automated Analysis of Privacy Policies with Large Language Models](#) (September 2023)
- [LLM4Jobs: Unsupervised occupation extraction and standardization leveraging Large Language Models](#) (September 2023)
- [Summarization is \(Almost\) Dead](#) (September 2023)
- [Investigating Zero- and Few-shot Generalization in Fact Verification](#) (September 2023)
- [Performance of the Pre-Trained Large Language Model GPT-4 on Automated Short Answer Grading](#) (September 2023)
- [Contrastive Decoding Improves Reasoning in Large Language Models](#) (September 2023)
- [Struc-Bench: Are Large Language Models Really Good at Generating Complex Structured Data?](#) (September 2023)
- [Neural Machine Translation Models Can Learn to be Few-shot Learners](#) (September 2023)
- [Chain-of-Thought Reasoning is a Policy Improvement Operator](#) (September 2023)
- [ICLEF: In-Context Learning with Expert Feedback for Explainable Style Transfer](#) (September 2023)
- [When do Generative Query and Document Expansions Fail? A Comprehensive Study Across Methods, Retrievers, and Datasets](#) (September 2023)
- [Using Large Language Models for Knowledge Engineering \(LLMKE\): A Case Study on Wikidata](#) (September 2023)
- [Self-Consistent Narrative Prompts on Abductive Natural Language Inference](#) (September 2023)
- [Investigating Answerability of LLMs for Long-Form Question Answering](#) (September 2023)
- [PromptTTS++: Controlling Speaker Identity in Prompt-Based Text-to-Speech Using Natural Language Descriptions](#) (September 2023)

- [An Empirical Evaluation of Prompting Strategies for Large Language Models in Zero-Shot Clinical Natural Language Processing](#) (September 2023)
- [Leveraging Contextual Information for Effective Entity Salience Detection](#) (September 2023)
- [Prompting4Debugging: Red-Teaming Text-to-Image Diffusion Models by Finding Problematic Prompts](#) (September 2023)
- [PACE: Prompting and Augmentation for Calibrated Confidence Estimation with GPT-4 in Cloud Incident Root Cause Analysis](#) (September 2023)
- [From Sparse to Dense: GPT-4 Summarization with Chain of Density Prompting](#) (September 2023)
- [Measuring and Improving Chain-of-Thought Reasoning in Vision-Language Models](#) (September 2023)
- [Zero-Resource Hallucination Prevention for Large Language Models](#) (September 2023)
- [Certifying LLM Safety against Adversarial Prompting](#) (September 2023)
- [Improving Code Generation by Dynamic Temperature Sampling](#) (September 2023)
- [Prompting a Large Language Model to Generate Diverse Motivational Messages: A Comparison with Human-Written Messages](#) (August 2023)
- [Financial News Analytics Using Fine-Tuned Llama 2 GPT Model](#) (August 2023)
- [A Study on Robustness and Reliability of Large Language Model Code Generation](#) (August 2023)
- [Large Language Models Vote: Prompting for Rare Disease Identification](#) (August 2023)
- [WizardMath: Empowering Mathematical Reasoning for Large Language Models via Reinforced Evol-Instruct](#) (August 2023)
- [Tree-of-Mixed-Thought: Combining Fast and Slow Thinking for Multi-hop Visual Reasoning](#) (August 2023)
- [Graph of Thoughts: Solving Elaborate Problems with Large Language Models](#) (August 2023)
- [Red-Teaming Large Language Models using Chain of Utterances for Safety-Alignment](#) (August 2023)
- [Boosting Logical Reasoning in Large Language Models through a New Framework: The Graph of Thought](#) (August 2023)
- [You Only Prompt Once: On the Capabilities of Prompt Learning on Large Language Models to Tackle Toxic Content](#) (August 2023)
- [LLM As DBA](#) (August 2023)
- [Interpretable Math Word Problem Solution Generation Via Step-by-step Planning](#) (June 2023)
- [In-Context Learning User Simulators for Task-Oriented Dialog Systems](#) (June 2023)

- [SQL-PaLM: Improved Large Language Model Adaptation for Text-to-SQL](#) (June 2023)
- [Effective Structured Prompting by Meta-Learning and Representative Verbalizer](#) (June 2023)
- [Layout and Task Aware Instruction Prompt for Zero-shot Document Image Question Answering](#) (June 2023)
- [Chain-Of-Thought Prompting Under Streaming Batch: A Case Study](#) (June 2023)
- [Red Teaming Language Model Detectors with Language Models](#) (May 2023)
- [Gorilla: Large Language Model Connected with Massive APIs](#) (May 2023)
- [Deliberate then Generate: Enhanced Prompting Framework for Text Generation](#) (May 2023)
- [What does the Failure to Reason with "Respectively" in Zero/Few-Shot Settings Tell Us about Language Models?](#) (May 2023)
- [ScoNe: Benchmarking Negation Reasoning in Language Models With Fine-Tuning and In-Context Learning](#) (May 2023)
- [SheetCopilot: Bringing Software Productivity to the Next Level through Large Language Models](#) (May 2023)
- [Grammar Prompting for Domain-Specific Language Generation with Large Language Models](#) (May 2023)
- [Mitigating Label Biases for In-context Learning](#) (May 2023)
- [Short Answer Grading Using One-shot Prompting and Text Similarity Scoring Model](#) (May 2023)
- [Strategic Reasoning with Language Models](#) (May 2023)
- [Dissecting Chain-of-Thought: A Study on Compositional In-Context Learning of MLPs](#) (May 2023)
- [Marked Personas: Using Natural Language Prompts to Measure Stereotypes in Language Models](#) (May 2023)
- [Leveraging Training Data in Few-Shot Prompting for Numerical Reasoning](#) (May 2023)
- [Exploring Effectiveness of GPT-3 in Grammatical Error Correction: A Study on Performance and Controllability in Prompt-Based Methods](#) (May 2023)
- [NOTABLE: Transferable Backdoor Attacks Against Prompt-based NLP Models](#) (May 2023)
- [Tab-CoT: Zero-shot Tabular Chain of Thought](#) (May 2023)
- [Evaluating GPT-3 Generated Explanations for Hateful Content Moderation](#) (May 2023)
- [Prompt-Guided Retrieval Augmentation for Non-Knowledge-Intensive Tasks](#) (May 2023)
- [Zero- and Few-Shot Event Detection via Prompt-Based Meta Learning]<https://arxiv.org/abs/2305.17373> (May 2023)

- [Chain-of-Thought Hub: A Continuous Effort to Measure Large Language Models' Reasoning Performance](#) (May 2023)
- [Large Language Models Can be Lazy Learners: Analyze Shortcuts in In-Context Learning](#) (May 2023)
- [Heterogeneous Value Evaluation for Large Language Models](#) (May 2023)
- [PromptNER: Prompt Locating and Typing for Named Entity Recognition](#) (May 2023)
- [Small Language Models Improve Giants by Rewriting Their Outputs](#) (May 2023)
- [On the Planning Abilities of Large Language Models -- A Critical Investigation](#) (May 2023)
- [Beyond Chain-of-Thought, Effective Graph-of-Thought Reasoning in Large Language Models](#) (May 2023)
- [PRODIGY: Enabling In-context Learning Over Graphs](#) (May 2023)
- [Large Language Models are Few-Shot Health Learners](#) (May 2023)
- [Role-Play with Large Language Models](#) (May 2023)
- [Measuring Inductive Biases of In-Context Learning with Underspecified Demonstrations](#) (May 2023)
- [Fact-Checking Complex Claims with Program-Guided Reasoning](#) (May 2023)
- [Large Language Models as Tool Makers](#) (May 2023)
- [Iterative Forward Tuning Boosts In-context Learning in Language Models](#) (May 2023)
- [SwiftSage: A Generative Agent with Fast and Slow Thinking for Complex Interactive Tasks](#) (May 2023)
- [Interactive Natural Language Processing](#) (May 2023)
- [An automatically discovered chain-of-thought prompt generalizes to novel models and datasets](#) (May 2023)
- [Large Language Model Guided Tree-of-Thought](#) (May 2023)
- [Active Retrieval Augmented Generation](#) (May 2023)
- [A PhD Student's Perspective on Research in NLP in the Era of Very Large Language Models](#) (May 2023)
- [Visual Chain of Thought: Bridging Logical Gaps with Multimodal Infillings](#) (May 2023)
- [Mirages: On Anthropomorphism in Dialogue Systems](#) (May 2023)
- [Model evaluation for extreme risks](#) (May 2023)
- [Language Models Don't Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting](#) (May 2023)
- [Cognitive Reframing of Negative Thoughts through Human-Language Model Interaction](#) (May 2023)
- [PromptClass: Weakly-Supervised Text Classification with Prompting Enhanced Noise-Robust Self-Training](#) (May 2023)

- [Augmented Large Language Models with Parametric Knowledge Guiding](#) (May 2023)
- [Aligning Large Language Models through Synthetic Feedback](#) (May 2023)
- [Concept-aware Training Improves In-context Learning Ability of Language Models](#) (May 2023)
- [FrugalGPT: How to Use Large Language Models While Reducing Cost and Improving Performance](#) (May 2023)
- [Enhancing Black-Box Few-Shot Text Classification with Prompt-Based Data Augmentation](#) (May 2023)
- [Detecting automatically the layout of clinical documents to enhance the performances of downstream natural language processing](#) (May 2023)
- ["Is the Pope Catholic?" Applying Chain-of-Thought Reasoning to Understanding Conversational Implicatures](#) (May 2023)
- [Let's Think Frame by Frame: Evaluating Video Chain of Thought with Video Infilling and Prediction](#) (May 2023)
- [Generating Data for Symbolic Language with Large Language Models](#) (May 2023)
- [Make a Choice! Knowledge Base Question Answering with In-Context Learning](#) (May 2023)
- [Improving Language Models via Plug-and-Play Retrieval Feedback](#) (May 2023)
- [Multi-Granularity Prompts for Topic Shift Detection in Dialogue](#) (May 2023)
- [The CoT Collection: Improving Zero-shot and Few-shot Learning of Language Models via Chain-of-Thought Fine-Tuning](#) (May 2023)
- [Can Language Models Understand Physical Concepts?](#) (May 2023)
- [Evaluating Factual Consistency of Summaries with Large Language Models](#) (May 2023)
- [Dr.ICL: Demonstration-Retrieved In-context Learning](#) (May 2023)
- [Probing in Context: Toward Building Robust Classifiers via Probing Large Language Models](#) (May 2023)
- [Skill-Based Few-Shot Selection for In-Context Learning](#) (May 2023)
- [Exploring Chain-of-Thought Style Prompting for Text-to-SQL](#) (May 2023)
- [Enhancing Chat Language Models by Scaling High-quality Instructional Conversations](#) (May 2023)
- [On Learning to Summarize with Large Language Models as References](#) (May 2023)
- [Learning to Generate Novel Scientific Directions with Contextualized Literature-based Discovery](#) (May 2023)
- [Active Learning Principles for In-Context Learning with Large Language Models](#) (May 2023)
- [Two Failures of Self-Consistency in the Multi-Step Reasoning of LLMs](#) (May 2023)

- [Improving Factuality and Reasoning in Language Models through Multiagent Debate](#) (May 2023)
- [ChatCoT: Tool-Augmented Chain-of-Thought Reasoning on Chat-based Large Language Models](#) (May 2023)
- [WikiChat: A Few-Shot LLM-Based Chatbot Grounded with Wikipedia](#) (May 2023)
- [Query Rewriting for Retrieval-Augmented Large Language Models](#) (May 2023)
- [Discrete Prompt Optimization via Constrained Generation for Zero-shot Re-ranker](#) (May 2023)
- [Element-aware Summarization with Large Language Models: Expert-aligned Evaluation and Chain-of-Thought Method](#) (May 2023)
- [Small Language Models Improve Giants by Rewriting Their Outputs](#) (May 2023)
- [Prompting and Evaluating Large Language Models for Proactive Dialogues: Clarification, Target-guided, and Non-collaboration](#) (May 2023)
- [Prompt-Based Monte-Carlo Tree Search for Goal-Oriented Dialogue Policy Planning](#) (May 2023)
- [Mitigating Language Model Hallucination with Interactive Question-Knowledge Alignment](#) (May 2023)
- [Making Language Models Better Tool Learners with Execution Feedback](#) (May 2023)
- [Text-to-SQL Error Correction with Language Models of Code](#) (May 2023)
- [Decomposed Prompting for Machine Translation Between Related Languages using Large Language Models](#) (May 2023)
- [SPARSEFIT: Few-shot Prompting with Sparse Fine-tuning for Jointly Generating Predictions and Natural Language Explanations](#) (May 2023)
- ["According to ..." Prompting Language Models Improves Quoting from Pre-Training Data](#) (May 2023)
- [Prompt-based methods may underestimate large language models' linguistic generalizations](#) (May 2023)
- [Chain of Knowledge: A Framework for Grounding Large Language Models with Structured Knowledge Bases](#) (May 2023)
- [Measuring Inductive Biases of In-Context Learning with Underspecified Demonstrations](#) (May 2023)
- [Automated Few-shot Classification with Instruction-Finetuned Language Models](#) (May 2023)
- [Enhancing Few-shot Text-to-SQL Capabilities of Large Language Models: A Study on Prompt Design Strategies](#) (May 2023)
- [MvP: Multi-view Prompting Improves Aspect Sentiment Tuple Prediction](#) (May 2023)
- [Learning Interpretable Style Embeddings via Prompting LLMs](#) (May 2023)
- [Enhancing Small Medical Learners with Privacy-preserving Contextual Prompting](#) (May 2023)

- [Fact-Checking Complex Claims with Program-Guided Reasoning](#) (May 2023)
- [A Benchmark on Extremely Weakly Supervised Text Classification: Reconcile Seed Matching and Prompting Approaches](#) (May 2023)
- [This Prompt is Measuring <MASK>: Evaluating Bias Evaluation in Language Models](#) (May 2023)
- [Enhancing Cross-lingual Natural Language Inference by Soft Prompting with Multilingual Verbalizer](#) (May 2023)
- [Evaluating Prompt-based Question Answering for Object Prediction in the Open Research Knowledge Graph](#) (May 2023)
- [Explaining How Transformers Use Context to Build Predictions](#) (May 2023)
- [PiVe: Prompting with Iterative Verification Improving Graph-based Generative Capability of LLMs](#) (May 2023)
- [PromptNER: A Prompting Method for Few-shot Named Entity Recognition via k Nearest Neighbor Search](#) (May 2023)
- [Logic-LM: Empowering Large Language Models with Symbolic Solvers for Faithful Logical Reasoning](#) (May 2023)
- [Enhancing Few-shot NER with Prompt Ordering based Data Augmentation](#) (May 2023)
- [Chain-of-thought prompting for responding to in-depth dialogue questions with LLM](#) (May 2023)
- [How to Prompt LLMs for Text-to-SQL: A Study in Zero-shot, Single-domain, and Cross-domain Settings](#) (May 2023)
- [Evaluation of medium-large Language Models at zero-shot closed book generative question answering](#) (May 2023)
- [Few-Shot Dialogue Summarization via Skeleton-Assisted Prompt Transfer](#) (May 2023)
- [Can NLP Models Correctly Reason Over Contexts that Break the Common Assumptions?](#) (May 2023)
- [Reasoning Implicit Sentiment with Chain-of-Thought Prompting](#) (May 2023)
- [Writing your own book: A method for going from closed to open book QA to improve robustness and performance of smaller LLMs](#) (May 2023)
- [AutoTrial: Prompting Language Models for Clinical Trial Design](#) (May 2023)
- [CRITIC: Large Language Models Can Self-Correct with Tool-Interactive Critiquing](#) (May 2023)
- [Controlling the Extraction of Memorized Data from Large Language Models via Prompt-Tuning](#) (May 2023)
- [Prompting with Pseudo-Code Instructions](#) (May 2023)
- [TrueTeacher: Learning Factual Consistency Evaluation with Large Language Models](#) (May 2023)

- [Aligning Instruction Tasks Unlocks Large Language Models as Zero-Shot Relation Extractors](#) (May 2023)
- [Exploiting Biased Models to De-bias Text: A Gender-Fair Rewriting Model](#) (May 2023)
- [Learning In-context Learning for Named Entity Recognition](#) (May 2023)
- [Take a Break in the Middle: Investigating Subgoals towards Hierarchical Script Generation](#) (May 2023)
- [TEPrompt: Task Enlightenment Prompt Learning for Implicit Discourse Relation Recognition](#) (May 2023)
- [Large Language Models can be Guided to Evade AI-Generated Text Detection](#) (May 2023)
- [Temporal Knowledge Graph Forecasting Without Knowledge Using In-Context Learning](#) (May 2023)
- [Prompting the Hidden Talent of Web-Scale Speech Models for Zero-Shot Task Generalization](#) (May 2023)
- [Think Outside the Code: Brainstorming Boosts Large Language Models in Code Generation](#) (May 2023)
- [Improving Language Model Negotiation with Self-Play and In-Context Learning from AI Feedback](#) (May 2023)
- [ConvXAI: Delivering Heterogeneous AI Explanations via Conversations to Support Human-AI Scientific Writing](#) (May 2023)
- [StructGPT: A General Framework for Large Language Model to Reason over Structured Data](#) (May 2023)
- [Towards Expert-Level Medical Question Answering with Large Language Models](#) (May 2023)
- [Large Language Models are Built-in Autoregressive Search Engines](#) (May 2023)
- [MsPrompt: Multi-step Prompt Learning for Debiasing Few-shot Event Detection](#) (May 2023)
- [Exploring the Impact of Layer Normalization for Zero-shot Neural Machine Translation](#) (May 2023)
- [SGP-TOD: Building Task Bots Effortlessly via Schema-Guided LLM Prompting](#) (May 2023)
- [Multi-modal Visual Understanding with Prompts for Semantic Information Disentanglement of Image](#) (May 2023)
- [Soft Prompt Decoding for Multilingual Dense Retrieval](#) (May 2023)
- [PaLM 2 Technical Report](#) (May 2023)
- [Are LLMs All You Need for Task-Oriented Dialogue?](#) (April 2023)
- [HiPrompt: Few-Shot Biomedical Knowledge Fusion via Hierarchy-Oriented Prompting](#) (April 2023)
- [Approximating Human Evaluation of Social Chatbots with Prompting](#) (April 2023)

- [Automated Reading Passage Generation with OpenAI's Large Language Model](#) (April 2023)
- [WebBrain: Learning to Generate Factually Correct Articles for Queries by Grounding on Large Web Corpus](#) (April 2023)
- [Prompt Pre-Training with Twenty-Thousand Classes for Open-Vocabulary Visual Recognition](#) (April 2023)
- [GPT detectors are biased against non-native English writers](#) (April 2023)
- [Zero-Shot Next-Item Recommendation using Large Pretrained Language Models](#) (April 2023)
- [Large Language Models as Master Key: Unlocking the Secrets of Materials Science with GPT](#) (April 2023)
- [Efficiently Aligned Cross-Lingual Transfer Learning for Conversational Tasks using Prompt-Tuning](#) (April 2023)
- [Better Language Models of Code through Self-Improvement](#) (April 2023)
- [PromptORE -- A Novel Approach Towards Fully Unsupervised Relation Extraction](#) (April 2023)
- [Assessing Language Model Deployment with Risk Cards](#) (April 2023)
- [Enhancing Large Language Models with Climate Resources](#) (March 2023)
- [BloombergGPT: A Large Language Model for Finance](#) (March 2023)
- [Medical Intervention Duration Estimation Using Language-enhanced Transformer Encoder with Medical Prompts](#) (March 2023)
- [Soft-prompt tuning to predict lung cancer using primary care free-text Dutch medical notes](#) (March 2023)
- [TaskMatrix.AI: Completing Tasks by Connecting Foundation Models with Millions of APIs](#) (March 2023)
- [Larger Probes Tell a Different Story: Extending Psycholinguistic Datasets Via In-Context Learning](#) (March 2023)
- [Linguistically Informed ChatGPT Prompts to Enhance Japanese-Chinese Machine Translation: A Case Study on Attributive Clauses](#) (March 2023)
- [Knowledge-augmented Frame Semantic Parsing with Hybrid Prompt-tuning](#) (March 2023)
- [Debiasing Scores and Prompts of 2D Diffusion for Robust Text-to-3D Generation](#) (March 2023)
- [Zero-shot Model Diagnosis](#) (March 2023)
- [Prompting Large Language Models to Generate Code-Mixed Texts: The Case of South East Asian Languages](#) (March 2023)
- [SPeC: A Soft Prompt-Based Calibration on Mitigating Performance Variability in Clinical Notes Summarization](#) (March 2023)
- [Large Language Models and Simple, Stupid Bugs](#) (March 2023)

- [Can Generative Pre-trained Transformers \(GPT\) Pass Assessments in Higher Education Programming Courses?](#) (March 2023)
- [SelfCheckGPT: Zero-Resource Black-Box Hallucination Detection for Generative Large Language Models](#) (March 2023)
- [Large Language Models in the Workplace: A Case Study on Prompt Engineering for Job Type Classification](#) (March 2023)
- [ICL-D3IE: In-Context Learning with Diverse Demonstrations Updating for Document Information Extraction](#) (March 2023)
- [MathPrompter: Mathematical Reasoning using Large Language Models](#) (March 2023)
- [Prompt-Based Learning for Thread Structure Prediction in Cybersecurity Forums](#) (March 2023)
- [Choice Over Control: How Users Write with Large Language Models using Diegetic and Non-Diegetic Prompting](#) (March 2023)
- [Prompting Large Language Models with Answer Heuristics for Knowledge-based Visual Question Answering](#) (March 2023)
- [Soft Prompt Guided Joint Learning for Cross-Domain Sentiment Analysis](#) (March 2023)
- [SpeechPrompt v2: Prompt Tuning for Speech Classification Tasks](#) (March 2023)
- [Goal Driven Discovery of Distributional Differences via Language Descriptions](#) (February 2023)
- [Navigating the Grey Area: Expressions of Overconfidence and Uncertainty in Language Models](#) (February 2023)
- [TabGenie: A Toolkit for Table-to-Text Generation](#) (February 2023)
- [SGL-PT: A Strong Graph Learner with Graph Prompt Tuning](#) (February 2023)
- [Few-Shot Table-to-Text Generation with Prompt-based Adapter](#) (February 2023)
- [Language Models Are Few-shot Learners for Prognostic Prediction](#) (February 2023)
- [STA: Self-controlled Text Augmentation for Improving Text Classifications](#) (February 2023)
- [Check Your Facts and Try Again: Improving Large Language Models with External Knowledge and Automated Feedback](#) (February 2023)
- [How Generative AI models such as ChatGPT can be \(Mis\)Used in SPC Practice, Education, and Research? An Exploratory Study](#) (February 2023)
- [Grimm in Wonderland: Prompt Engineering with Midjourney to Illustrate Fairytales](#) (February 2023)
- [LabelPrompt: Effective Prompt-based Learning for Relation Classification](#) (February 2023)
- [Language Model Crossover: Variation through Few-Shot Prompting](#) (February 2023)

- [Prompt Tuning of Deep Neural Networks for Speaker-adaptive Visual Speech Recognition](#) (February 2023)
- [The Capacity for Moral Self-Correction in Large Language Models](#) (February 2023)
- [Prompting for Multimodal Hateful Meme Classification](#) (February 2023)
- [PLACES: Prompting Language Models for Social Conversation Synthesis](#) (February 2023)
- [Toolformer: Language Models Can Teach Themselves to Use Tools](#) (February 2023)
- [Commonsense-Aware Prompting for Controllable Empathetic Dialogue Generation](#) (February 2023)
- [Crawling the Internal Knowledge-Base of Language Models](#) (January 2023)
- [Legal Prompt Engineering for Multilingual Legal Judgement Prediction](#) (December 2022)
- [Investigating Prompt Engineering in Diffusion Models](#) (November 2022)
- [Learn to Explain: Multimodal Reasoning via Thought Chains for Science Question Answering](#) (September 2022)
- [Conversing with Copilot: Exploring Prompt Engineering for Solving CS1 Problems Using Natural Language](#) (October 2022)
- [Piloting Copilot and Codex: Hot Temperature, Cold Prompts, or Black Magic?](#) (October 2022)
- [Plot Writing From Scratch Pre-Trained Language Models](#) (July 2022)
- [Survey of Hallucination in Natural Language Generation](#) (February 2022)

Collections

- [Chain-of-Thought Papers](#)
- [Papers with Code](#)
- [Prompt Papers](#)

Last updated on September 19, 2024