Few-Shot Sentiment Classification with LLMs

Background

This prompt tests an LLM's text classification capabilities by prompting it to classify a piece of text into the proper sentiment using few-shot examples.

Prompt

```
This is awesome! // Negative
This is bad! // Positive
Wow that movie was rad! // Positive
What a horrible show! //
```

Code / API

GPT-4 (OpenAI)   Mixtral MoE 8x7B Instruct (Fireworks)

```python
from openai import OpenAI
client = OpenAI()

response = client.chat.completions.create(
    model="gpt-4",
    messages=[
        {
        "role": "user",
        "content": "This is awesome! // Negative\nThis is bad! // Positive\nWow that movie
was rad! // Positive\nWhat a horrible show! //"
        }
    ],
    temperature=1,
    max_tokens=256,
    top_p=1,
    frequency_penalty=0,
    presence_penalty=0
)
```

Reference

- [Prompt Engineering Guide](#) (16 March 2023)

Last updated on September 19, 2024