



# ECOAR – Computer Architecture

## Lecture notes

### Module 2 – Data

Electrical and Computer Engineering  
Faculty of Electronics and Information Technology  
Warsaw University of Technology



HUMAN CAPITAL  
HUMAN – BEST INVESTMENT!

EUROPEAN UNION  
EUROPEAN  
SOCIAL FUND



## Outline

- Data types
- Binary data representation
- Data in computer's memory
- Vector data types

# Data types

- Boolean (True/False)
- Text characters
- Numbers
  - integer
    - unsigned
    - signed
  - non-integer
    - Fixed point
    - Floating point
- Sounds and other unidimensional signals
- Raster images

# Binary data representation

- Computer operates solely on groups of bits (binary digits), treating them as binary numbers
  - Contemporary computers usually operate on binary words of lengths equal  $8 \times 2^n$  bits (8, 16, 32, 64)
- Non-numeric data must be expressed using numbers

# Alphanumeric data

- Every character is represented by a number denoting its position in code table
- Most frequently used codes:
  - ASCII – 128 positions including upper- and lowercase Latin letters
  - Extended ASCII – 256 positions
    - First 128 positions identical to ASCII, the remaining 128 positions used for national characters and special symbols
    - problem: different codes are needed for various parts of the world
  - EBCDIC family of codes – 256 symbols, used mainly by IBM
  - UNICODE
    - Initially  $2^{16}$ , currently up to  $2^{21}$  positions
    - Currently contains approx. 150000 characters
    - Covers all alphabetic characters used in the world

# ASCII

- American Standard Code for Information Interchange
- Developed for teletype devices, later used in computers
- 128 positions, 95 visible and 33 invisible
- Invisible – use codes 0..31 (0x00..0x1f) and 127 (0x7f):
  - White spaces and formatting codes (tabs, new line, new page)
  - Transmission and device control
  - CAN – cancel/delete - code 127 (0x7f)
- Visible – use codes 32..126 (0x20..0x7e):
  - space (code 32)
  - Digits
  - Latin letters
  - Punctuation marks
  - Common mathematical symbols

# What a programmer should know about ASCII

- Control codes - positions 0..31, including
  - CR – carriage return – 13
  - LF – line feed – 10
  - Also important: HT, FF, BSP, BEL
- Space – 32 (0x20)
- digits 0..9 – codes 48..57 (0x30..0x39)
  - to get digit's value, subtract `0` (0x30) from character representing it
- Letters in alphabetic order
  - Upper case – 65..90 (0x41..0x5a)
  - Lower case – 97..122 (0x61..0x7a)
  - Distance between lower and upper case of a letter is 32 (0x20)
- Other visible chars fit in the range 33..126
- 127 – the very last special code (cancel/delete)

## Extended codes based on ASCII

- 256 code points – 8-bit representation
  - First 128 positions identical to ASCII
  - Next 128 positions used to represent characters from a selected set of alphabets, ex.
    - Slavic
    - Nordic
    - Cyrillic
    - Greek
- Many code tables of different origin used in some countries
  - ISO8859 family (Polish characters present in ISO8859-2)
  - Microsoft – “code pages” - CP (Poland – CP852, CP1250)
  - Local usage (Polish - Mazovia, Polgaz)



# UNICODE and UTF-8

- Unicode, introduced in late 1980s, is a 21-bit character code designed to include all alphabetic characters used in the world
- Symbolic notation of Unicode character: U+<hex\_code>
- UTF-8 is the most common Unicode representation
- A character is represented by 1..4 octets  
character binary code

b <sub>20</sub>	b <sub>19</sub>	b <sub>18</sub>	b <sub>17</sub>	b <sub>16</sub>	b <sub>15</sub>	b <sub>14</sub>	b <sub>13</sub>	b <sub>13</sub>	b <sub>12</sub>	b <sub>11</sub>	b <sub>10</sub>	b <sub>9</sub>	b <sub>8</sub>	b <sub>7</sub>	b <sub>6</sub>	b <sub>5</sub>	b <sub>4</sub>	b <sub>3</sub>	b <sub>2</sub>	b <sub>1</sub>	b <sub>0</sub>
-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	-----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------	----------------

is represented as:

- ASCII chars U+00..U+7F → 1 octet: 0b<sub>6</sub>b<sub>5</sub>b<sub>4</sub>b<sub>3</sub>b<sub>2</sub>b<sub>1</sub>b<sub>0</sub>
- U+80..U+7FF → 2 octets: 110b<sub>10</sub>b<sub>9</sub>b<sub>8</sub>b<sub>7</sub>b<sub>6</sub> 10b<sub>5</sub>b<sub>4</sub>b<sub>3</sub>b<sub>2</sub>b<sub>1</sub>b<sub>0</sub>
- U+800..U+FFFF → 3 octets:  
1110b<sub>15</sub>b<sub>14</sub>b<sub>13</sub>b<sub>12</sub> 10b<sub>11</sub>b<sub>10</sub>b<sub>9</sub>b<sub>8</sub>b<sub>7</sub>b<sub>6</sub> 10b<sub>5</sub>b<sub>4</sub>b<sub>3</sub>b<sub>2</sub>b<sub>1</sub>b<sub>0</sub>
- U+10000..U+10FFFF → 4 octets:  
11110b<sub>20</sub>b<sub>19</sub>b<sub>18</sub> 10b<sub>17</sub>b<sub>16</sub>b<sub>15</sub>b<sub>14</sub>b<sub>13</sub>b<sub>12</sub> 10b<sub>11</sub>b<sub>10</sub>b<sub>9</sub>b<sub>8</sub>b<sub>7</sub>b<sub>6</sub> 10b<sub>5</sub>b<sub>4</sub>b<sub>3</sub>b<sub>2</sub>b<sub>1</sub>b<sub>0</sub>



# Text strings

- Text string is a sequence of characters
- The software must know where the string ends
- Two common conventions:
  - Character string end identified by a special character
    - In C and similar languages, the string ends with NUL (zero) code
    - in other environments '\$' or LF was used
  - An explicit string length specified as binary number, stored before the text
    - common in Basic language

# Sounds and images

- Sound:
  - Voltage analog of acoustic pressure sampled with frequency dependent on the quality requirements (from 8 to 48 kHz)
  - Sample values stored as integer numbers
- Raster images
  - Represented as rectangular arrays of square picture elements (pixels)
  - Single color is assigned to every pixel
  - Color is represented using three primaries – values of basic lights (red, green, blue)
  - Values of primaries stored as unsigned integer numbers

# Units

- bit (Binary digIT) - abbreviation „b”- smallest information unit, may represent two-state information YES-NO, 1-0, TRUE-FALSE
- octet (abbrev. ‘o’) – 8 bits, tetrad/nibble – 4 bits
- byte – abbreviation „B” - smallest unit of information addressed by computer – in contemporary g.p computers – 8 bits
- word – unit of information operated on by the computer
  - 1, 2, 4, 8, 16 bytes
- Processor word – unit of information on which a given processor may operate
  - Maximum word size equals register width – in contemporary g.p. computers 64 or 32 bits
- Memory word – unit of information that may be transferred in a single transfer cycle between processor and memory
  - Typ. 128 or 64 bits
  - Longer memory word = faster data transfer

# Data formats

- Computers operate on words – groups of (usually)  $2^n$  bits
  - Common word lengths in g.p. computers – 8, 16, 32, 64, 128 bits
  - Specialized computers may use different word lengths – 24 bits
- Some computers are also capable of operating on single bits and bit fields
- Data stored as words in memory

## Logical (Boolean) data

- Single bit is sufficient to store Boolean data
  - Single bit Boolean format is used in computers providing hardware handling of single bit data
- Boolean data is commonly represented using word of a length natural for a given computer
  - `_Bool/bool` type in C/C++ is implemented as byte
  - Some other languages use the same size as used for integers – typ. 32 bits
- Representation depends on the programming environment and programming language
  - FALSE is usually represented by a word of value 0
  - TRUE may be represented by
    - Value 1 (C language – result)
    - Any non-zero value (most languages – source arguments)
    - Word made of all bits set to 1 (Visual Basic – result)

# Unsigned integers



- Natural binary code

$$v_N = \sum_{i=0}^{n-1} b_i \cdot 2^i$$

- BCD (Binary-Coded Decimal)

- Used for fixed point decimal numbers
- Decimal digits encoded in binary – 4 bits (nibble/tetrad) per digit
- Allowed nibble values 0..9
- Formats:
  - packed – 2 digits per byte; range 0..99
  - Unpacked (“ASCII”) - one digit per byte, range 0..9

# Signed integers

- Two's complement

$$v_{2C} = -b_{n-1} \cdot 2^{n-1} + \sum_{i=0}^{n-2} b_i \cdot 2^i$$

- Ones' complement

$$v_{1C} = -b_{n-1} \cdot (2^{n-1} - 1) + \sum_{i=0}^{n-2} b_i \cdot 2^i$$

- Sign-magnitude

$$v_{S-M} = (-1)^{b_{n-1}} \cdot \sum_{i=0}^{n-2} b_i \cdot 2^i$$

- Biased

- typically  
BIAS =  $2^{n-1} - 1$

$$v_B = -BIAS + \sum_{i=0}^{n-1} b_i \cdot 2^i$$



## Important features of numeric codes

- Representation of zero
  - Number of distinct representations
    - Two representations of 0 in ones complement and sign-magnitude
  - Ease of detection
- Symmetry of range (for signed numbers)
- Representation of sign
  - Ease of sign determination
- Change of sign
  - Ones complement – bit negation
  - Two's complement – bit negation and increment
  - Sign-magnitude – negation of sign bit
- Implementation of arithmetic operations
  - Addition and subtraction of two's complement numbers may be performed using the hardware designed for unsigned numbers
    - Only the overflow detection is different

# Numeric codes and their features

Code	Binary patterns			8-bit byte value		16-bit word value		32-bit word value	
	min	zero	max	min	max	min	max	min	max
Unsigned	00...0	00...0	11...1	0	255	0	65535	0	4294967295
2C	10...0	00...0	01...1	-128	127	-32768	32767	-2147483648	2147483647
S-M	11...1	00...0 10...0	01...1	-127	127	-32767	32767	-2147483647	2147483647
1C	10...0	00...0 11...1	01...1	-127	127	-32767	32767	-2147483647	2147483647
Biased	00...0	01...1	11...1	-127	128	-32767	32768	-2147483647	2147483648


# Change of sign, 2's complement dec→bin conversion

- Change of sign:
  - Ones' complement:  $\sim x$
  - Two's complement:  $\sim x + 1$
  - Biased:  $\text{BIAS} - x$
  - Sign-magnitude: negate the sign bit only
  - In codes with asymmetric range (2's C, biased), sign change may cause overflow
- Decimal to 2's complement conversion
  - For non-negative numbers, 2's complement representation is the same as unsigned
    - The representable range limit is smaller –  $2^{n-1} - 1$
  - For negative numbers:
    - Convert  $|x|$  to unsigned binary
    - Change the sign: negate all bits, then add 1

# Fixed point notation

- Obtained by shifting bit weights in integer notation
  - Equivalent to multiplication of integer value by  $2^{-f}$  ( $f$  – no. of bits in fractional part)
- Used for 2's complement and unsigned numbers
- Commonly used formats:
  - One or two bits of integer part, rest of word used as fractional part
  - Half of the word used as integer part, the other half as fraction
- Arithmetic operations similar to integer operations
  - Scaling (shifting) required during multiplication and division
  - Do not require special instructions or hardware structures

# Fixed point notation - examples



Integer weights, 32-bit word



Fixed point weights, 16 bits in integer part, 16 bits in fractional part



Fixed point weights, 1 bit of integer part, 31 bits in fractional part



## Binary vs. decimal fractions - accuracy

- In a fractional part of a binary number, bit weights are  $2^{-i}$ 
  - $\frac{1}{2}$ ,  $\frac{1}{4}$ ,  $\frac{1}{8}$ ,  $\frac{1}{16}$ ,  $\frac{1}{32}$ , etc...
- As 10 is not equal to any power of 2, decimal fractions cannot be in general precisely represented as binary fractions with finite number of digits
  - We face the same problem in decimal system:  
 $\frac{1}{3} = 0.33333333333333333333333333333333...$
  - Only fractions being sums of binary fractions may be represented precisely:  
 $0.75 = \frac{3}{4} = (0.11)_2$   
 $0.625 = \frac{5}{8} = (0.101)_2$   
...

## Floating point notation – basics

- Examples of decimal FP notation  
 $-1.234 \cdot 10^5$     $-0.1234 \cdot 10^6$     $-12.34 \cdot 10^4$
- Elements:
  - Sign of a number
  - Significand (unsigned fixed point number)
  - Exponent (signed integer)
- System base (10) is fixed – there is no need to store it
  - It is only a graphic element of the notation
- Normalized form – form in which integer part of significand is expressed using single non-zero digit
  - In the above example:  $-1.234 \cdot 10^5$
  - Other forms are denormalized (non normalized)
  - It is not possible to express zero in a normalized form

# Binary floating point notation – IEEE754, IEC60559

- If possible, numbers are stored in normalized form
  - Exceptions: zero and numbers of a very small magnitude
  - The only binary digit different from 0 is 1
    - Every normalized number has an integer part of significand equal to 1
    - A non-normalized number has an integer part of significand equal to 0
    - There is no need to store integer part of significand if we know whether the number is normalized or not
- Elements of notation
  - Sign – single bit: 0 – non-negative, 1 – non-positive
  - Exponent field
    - In normalized form – exponent in biased code
    - Special values: 00...00 – denormalized form, 11...11 – Not-a-Number
  - Trailing significand (mantissa) field – contains only fractional part of significand
    - Normalized form – integer part is 1
    - denormalized form – integer part is 0



# IEEE 754/IEC60559 – exponent field encoding

- exponent stored in e-bit wide field as biased value
- bias =  $2^{e-1} - 1$  (01..11 binary)

exp field	exponent value	meaning
00...00	-bias+1	zero or denormal
00...01	exp-bias = -bias+1	normalized value
0x..xx	exp-bias	normalized value
01..11	exp-bias = 0	normalized value
1x..xx	exp-bias	normalized value
11..10	exp-bias	biggest representable number
11..11	(none)	Not-a-Number

# IEEE754/IEC60559 - values

Sign	Biased exponent	Trailing significand (mantissa) field
S	E	T

Sign bit	Exponent field	Exp.value	Mantissa field	Significand value
x	00...00	-bias+1	m	0.m
x	from 00...01 to 11..10	exp-bias	m	1.m
0	11..11	-	00..00	+ infinity
1	11..11	-	00..00	- infinity
x	11..11	-	0x..xx	Signaling NaN
x	11..11	-	1x..xx	Quiet NaN

# IEEE754/IEC60559 – common formats

Word size [b]	Symbolic notation	Field widths		Application
		mantissa	exponent	
16	s10e5	10	5	Computer graphics data storage
24	s16e7	16	7	Computer graphics
16	s7e8	7	8	artificial intelligence (BF16)
<b>32</b>	<b>s23e8</b>	<b>23</b>	<b>8</b>	<b>binary32 (IEEE single), general purpose</b>
<b>64</b>	<b>s52e11</b>	<b>52</b>	<b>11</b>	<b>binary64 (IEEE double), general purpose</b>
80	s64e15	64	15	Extended double – x87
96	s80e15	80	15	DSP
128	s112e15	112	15	binary128, high-precision arithmetics

# Floating point arithmetic

- Floating point representation is an approximate one
  - Results of arithmetic operations are also approximate
- Result may depend on order of operations
  - $a + b + c$  may not be the same as  $c + b + a$
  - if  $|a| \ll |b|$ , then  $a + b$  may be equal to  $b$
  - Computing the sum or average of many values is not a trivial task
  - Addition and subtraction of multiple arguments should be performed in order of growing magnitude
- Equality test usually gives false result
  - Use  $\text{abs}(a-b) < \epsilon$  instead
- Precision of IEEE single (24 bits) is smaller than that of 32-bit integer or fixed point
  - Precise representation of integers in IEEE single is limited to absolute value range of  $2^{24}$  only

# Memory organization

- In general purpose computers 8-bit byte is the smallest addressable unit of memory
- Multibyte data occupy the appropriate number of consecutive byte-sized memory locations


## Multibyte data addressing – byte order

- Little Endian – least significant byte of a data word at the lowest address, more significant bytes follow
- Big Endian – most significant byte at the lowest address


32-bit word



Little-Endian



Big-Endian



# Little-Endian

- Byte numbering corresponds to bit weights
- Natural for computers
- Not so natural for humans
- Type casting preserves pointer value

32-bit word, value = 5

0	0	0	5
---	---	---	---

32-bit access

Addr+0:	5
Addr+1:	0
Addr+2:	0
Addr+3:	0

16-bit access

Addr+0:	5
Addr+1:	0

8-bit access

Addr+0:	5
---------	---

# Big-Endian

- Natural for humans
- Type cast changes the pointer value
- Fast string compare is possible

“abc” string in memory

Addr+0:	'a'
Addr+1:	'b'
Addr+2:	'c'
Addr+3:	0


32-bit word read from Addr

'a'	'b'	'c'	0
-----	-----	-----	---




## Logical vs. physical memory organization

- Logical – vector of bytes



- Physical - “2-dimensional” - bytes grouped in words, usually 32-, 64- or 128-bit
- Bytes within a word may be accessed simultaneously



# Data alignment

- Physically the memory is organized as a vector of words, each word being a vector of bytes
  - Any number of bytes within a single word may be accessed at the same time
- Memory word is usually at least as long as the processor's word
  - Typically 64 or 128 bits in 32-/64-bit processors (2 or 4 times longer than processor's word)
  - In older designs (before 1990) memory was sometimes narrower than processor
- The access to multibyte data contained in a single memory word is faster than in case when the data occupies parts of two memory words
  - In the latter case two physical accesses are needed
- New architectures enforce placement of every data item ensuring the fastest possible access

# Size alignment

- For the fastest access regardless of the physical memory width, each data item should be placed in memory at the address divisible by its length (rounded up to  $2^n$ )
- Size alignment is enforced by hardware in newer architectures
  - Unaligned access attempt results in execution error
- Even if not enforced in hardware, size alignment boosts the efficiency
  - Compilers enforce size alignment even if underlying hardware doesn't
- Each data type is characterized by two values:
  - Alignment – the data starting address must be a multiple of alignment
    - In C language returned by `_Alignof()` operator
  - Size – must be a multiple of alignment
    - In C language returned by `sizeof()` operator
    - in modern environments, size is always  $2^n$  and it is equal to alignment

# Vectors and arrays

- The elements of a vector occupy the consecutive memory locations, starting with the first element (placed at the lowest address – the address of a vector)
  - In C-like languages, the first element has an index of 0
    - $v[N]$  vector elements are  $v[0]$ ,  $v[1]$ , ...,  $v[N-1]$
- Vector alignment is the same as its element alignment
- An  $n$ -dimensional array is a vector of  $(n-1)$ -dimensional arrays

# Vectors and arrays

char s[12];      char t[3][4];      int32\_t v[12];      int32\_t a[3][4];


0	s[0]	0	t[0][0]	0	v[0]	0	a[0][0]
1	s[1]	1	t[0][1]	4	v[1]	4	a[0][1]
2	s[2]	2	t[0][2]	8	v[2]	8	a[0][2]
3	s[3]	3	t[0][3]	12	v[3]	12	a[0][3]
4	s[4]	4	t[1][0]	16	v[4]	16	a[1][0]
5	s[5]	5	t[1][1]	20	v[5]	20	a[1][1]
6	s[6]	6	t[1][2]	24	v[6]	24	a[1][2]
7	s[7]	7	t[1][3]	28	v[7]	28	a[1][3]
8	s[8]	8	t[2][0]	32	v[8]	32	a[2][0]
9	s[9]	9	t[2][1]	36	v[9]	36	a[2][1]
10	s[10]	10	t[2][2]	40	v[10]	40	a[2][2]
11	s[11]	11	t[2][3]	44	v[11]	44	a[2][3]
12	...	12	...	48	...	48	...

# Structures

- Compiler is required to preserve the order of fields
  - It cannot optimize the layout
- Each field must be aligned according to its type requirements
  - Padding (unused space) may be added between fields to achieve the proper alignment
- Structure must be aligned according to the alignment requirement of the field with the biggest alignment
- The structure is padded to the multiple of its alignment
- C language `sizeof()` operator returns the address distance between two structures of a given type, necessary for allocating a vector of structures.


# Order of structure fields vs. memory footprint

```
struct st1 {
    char c1;
    double d1;
    short int s1;
    double d2;
    char c2;
    int i1;
    char c3;
    short int s2;
    char c4;
    int i2;
} e1;
```



sizeof(struct st1) == 56

```
struct st2 {
    double d1, d2;
    int i1, i2;
    short int s1, s2;
    char c1, c2, c3, c4;
} e2;
```



sizeof(struct st2) == 32

## Vector data formats

- Contemporary processors may operate on long data words
- Sounds and pictures are expressed using short numbers
- Signal processing algorithms execute the same set of operations on all data samples
- It is possible to treat the long word as a vector of several shorter values
- The introduction of vector data formats enables better utilization of processor's resources



## Vector data formats – x86 SSE

- SSE is a vector unit implemented in x86 processors from 1999
- Data word – 128 bits
- Data formats supported (SSE2 and newer versions, 2000→):
  - 16 8-bit integers
  - 8 16-bit integers
  - 4 32-bit integers
  - 2 64-bit integers
  - 4 32-bit IEEE single
  - 2 64-bit IEEE double
  - 8 16-bit s10e5 floats (SSE5)

## x86 AVX vector unit

- new vector unit of x86 processors, introduced in 2011 in “2<sup>nd</sup> generation Intel Core i” processors
- 16 256-bit registers
- FP data formats similar to those of SSE but with 256-bit vectors – 4×64-bit or 8×32-bit
- AVX2 version also supports 256-bit integer vectors – 32×8b, 16×16b, 8×32b, 4×64b
- New extension called AVX512 supports 512-bit vectors

## Requirements for the test

- Integer conversion between decimal and 2's complement, sign-magnitude
- Floating point conversion: decimal  $\leftrightarrow$  IEEE754 binary32
- Structure and vector layout (drawing, sizeof operator value)
- Big-endian/little-endian data placement in memory



# ECOAR – Computer Architecture

## Lecture notes

### End of module

Electrical and Computer Engineering  
Faculty of Electronics and Information Technology  
Warsaw University of Technology



HUMAN CAPITAL  
HUMAN – BEST INVESTMENT!

EUROPEAN UNION  
EUROPEAN  
SOCIAL FUND

